Predicting Important Factors of NYC's Traffic Accidents

EL07: Syntax Error

1. Francis Nathan (U2120951H)
2. Matthew Dinata (U2120122G)
3. Muhammad Afiq (U2122587K)
4. Nigel Loh (U2122947F)

Introduction

Our Dataset

New York City Traffic Accidents from January - August 2020 classified into categories, such as CRASH DATE and TIME, CONTRIBUTING FACTOR VEHICLE, VEHICLE TYPE CODE, LATITUDE, LONGITUDE, etc.

Objective

Predict which factor most contributes to the number of persons injured in the traffic accident. This will help the government and police department in New York City set up suitable measurements and actionables to prevent more accidents.

Table of Contents

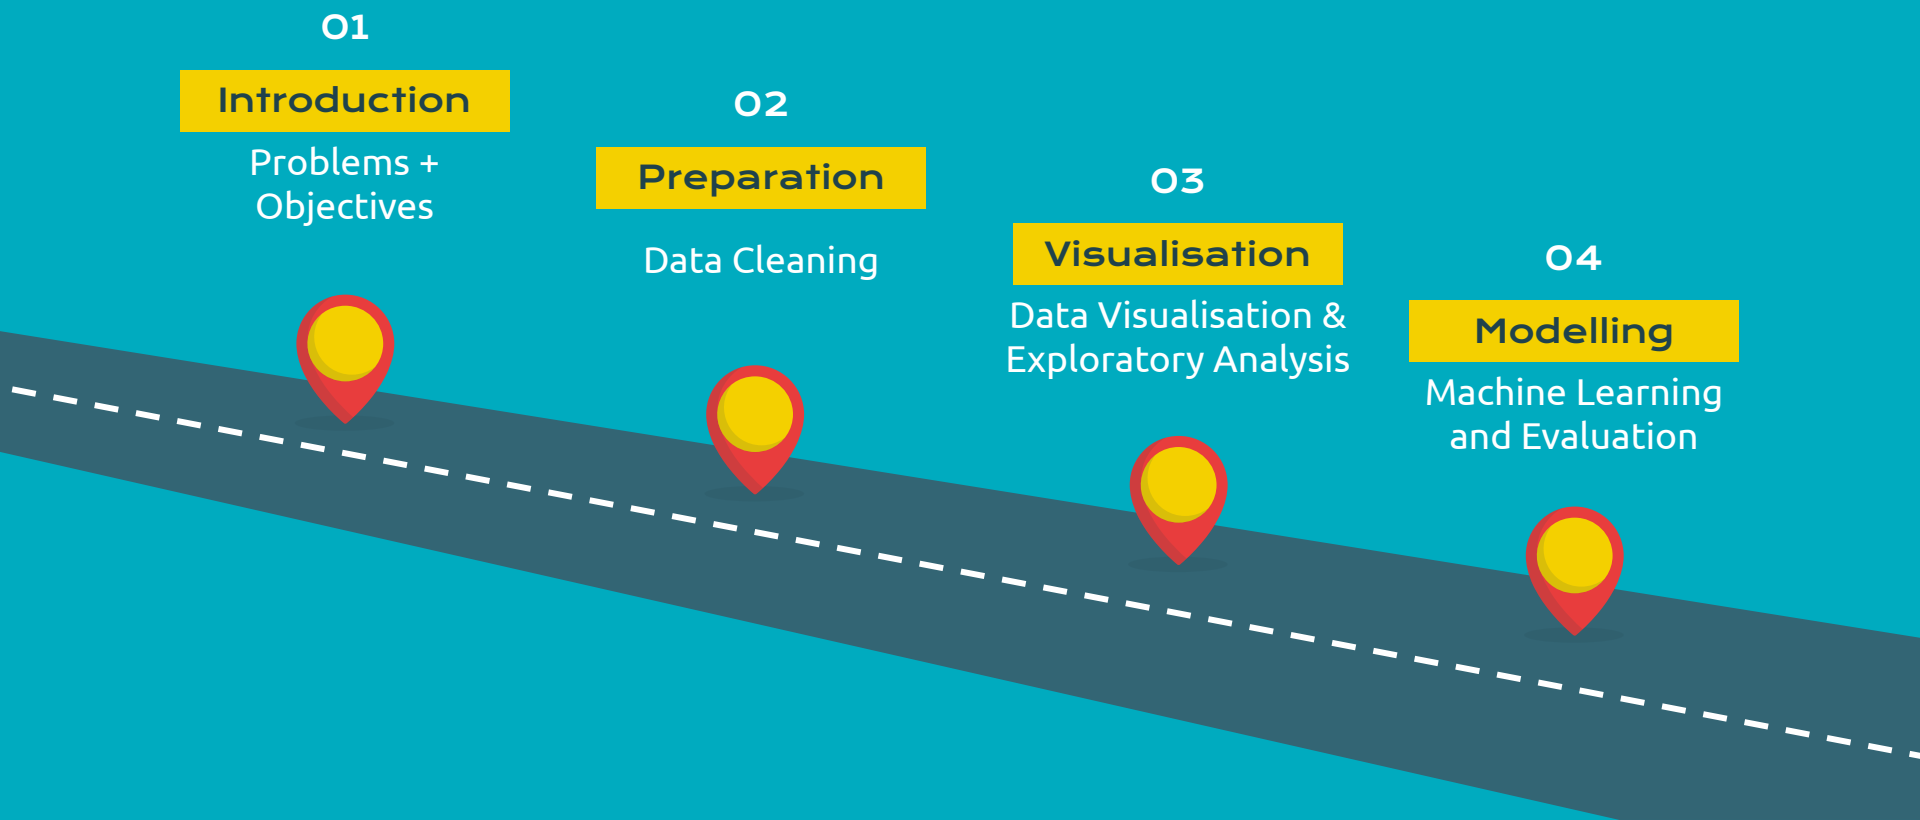

DATA PREPARATION

Data Preparation Contents

01

IMPORT

Importing packages and raw data

02

REMOVE COLUMNS

Dropping irrelevant columns

03

REMOVE NULL VALUES

Dropping NULL or ‘Unspecified values

04

DATA CONVERSION

Conversion of ‘CRASH DATE’ to Pandas DateTime

05

OVER-SAMPLING

Balancing data using RandomOverSampler

06

REDUCE DATA

Removal of random data rows to reduce data size

REMOVING NULL VALUES

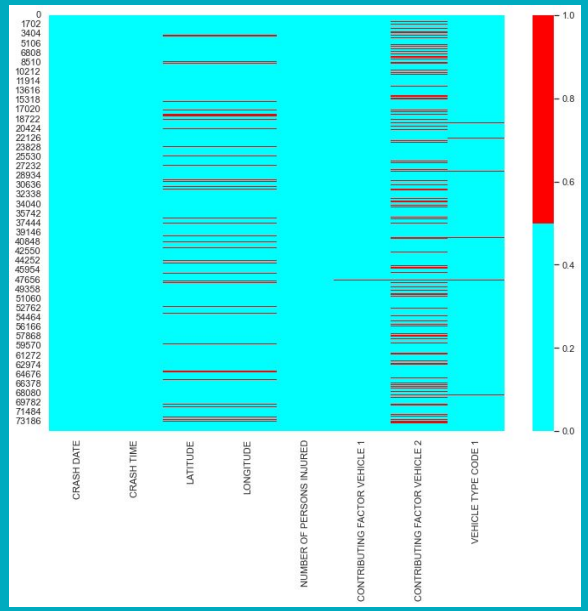

Based on the heatmap above, the color cyan indicates `No Missing Values` while the color red indicates `Missing Values`.

BALANCING DATA USING RANDOM OVER-SAMPLING

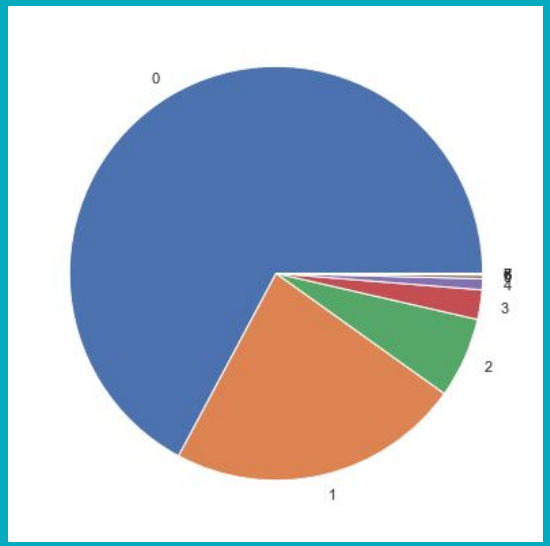

Comparison of the data before and after balancing using RandomOverSampler

REMOVAL OF RANDOM DATA ROWS

	CRASH DATE	CRASH TIME	LATITUDE	LONGITUDE	NUMBER OF PERSONS INJURED	CONTRIBUTING FACTOR VEHICLE 1	CONTRIBUTING FACTOR VEHICLE 2	VEHICLE TYPE CODE 1	VEHICLE TYPE CODE 2
0	2020-08-29	14:00:00	40.704422	-73.792854	0	Oversized Vehicle	Passing Too Closely	Bus	Station Wagon/Sport Utility Vehicle
1	2020-08-29	12:29:00	40.861862	-73.912820	2	Pavement Slippery	View Obstructed/Limited	Pick-up Truck	Station Wagon/Sport Utility Vehicle
3	2020-08-29	19:00:00	40.839680	-73.929276	1	Following Too Closely	Following Too Closely	Sedan	Station Wagon/Sport Utility Vehicle
4	2020-08-29	05:40:00	40.858190	-73.884350	0	Other Vehicular	Passing Too Closely	Sedan	Sedan
9	2020-08-29	15:00:00	40.669518	-73.911934	0	Other Vehicular	Other Vehicular	Station Wagon/Sport Utility Vehicle	Station Wagon/Sport Utility Vehicle
...	...	...	...	...	...	...	...	...	...
49528	2020-02-13	08:25:00	40.665230	-73.931465	8	Driver Inattention/Distraction	Driver Inattention/Distraction	Sedan	Station Wagon/Sport Utility Vehicle
49529	2020-02-14	08:40:00	40.854744	-73.923510	8	Driver Inattention/Distraction	Driver Inattention/Distraction	Sedan	Sedan
49530	2020-02-14	08:40:00	40.854744	-73.923510	8	Driver Inattention/Distraction	Driver Inattention/Distraction	Sedan	Sedan
49533	2020-02-13	08:25:00	40.665230	-73.931465	8	Driver Inattention/Distraction	Driver Inattention/Distraction	Sedan	Station Wagon/Sport Utility Vehicle
49534	2020-02-14	08:40:00	40.854744	-73.923510	8	Driver Inattention/Distraction	Driver Inattention/Distraction	Sedan	Sedan

24536 rows × 9 columns

Counter of 'NUMBER OF PERSONS INJURED':
{7: 2763, 2: 2753, 3: 2748, 4: 2748, 6: 2725, 0: 2722, 1: 2706, 5: 2699, 8: 2672}

DATA VISUALISATION

Data Visualisation Contents

01 LOCATION COORDINATES

02 CRASH DATE AND TIME

03 CONTRIBUTING FACTOR VEHICLE

04 VEHICLE TYPE CODE

LOCATION COORDINATES Visualisation

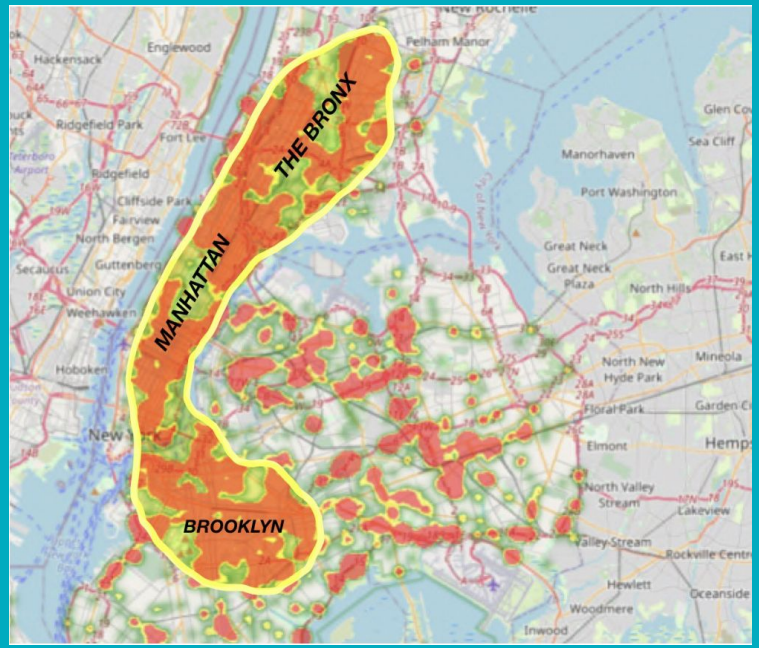

The map indicates that most of the traffic accidents in New York City occurs in **Brooklyn**, **Manhattan**, and **The Bronx.**

CONTRIBUTING VEHICLE TYPE Visualisation

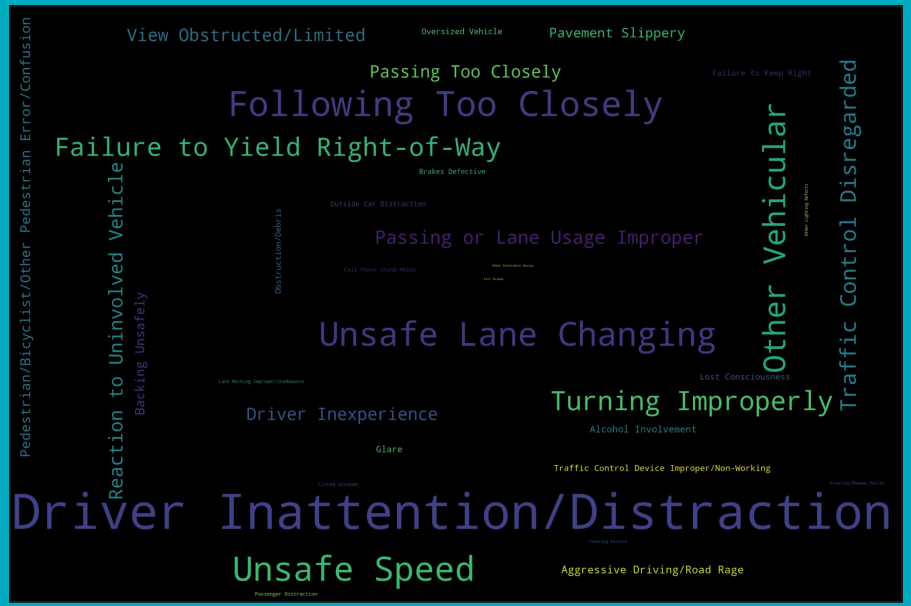

CONTRIBUTING VEHICLE TYPE 1 Word Cloud shows that most of the accidents are caused by Driver Inattention/Distraction

CONTRIBUTING VEHICLE TYPE 2 Word Cloud shows that most of the accidents are caused by Driver Inattention/Distraction

CRASH DATE and CRASH TIME Visualisation

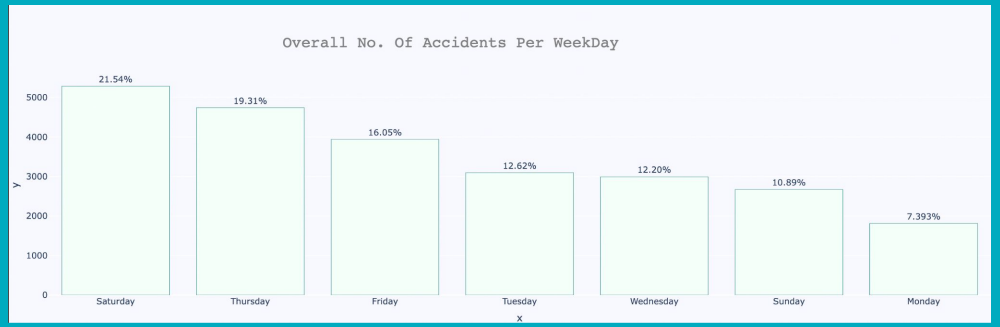

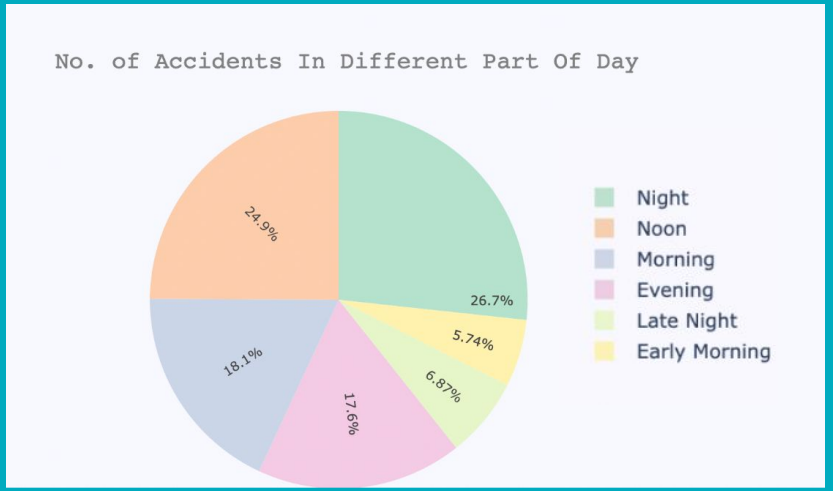

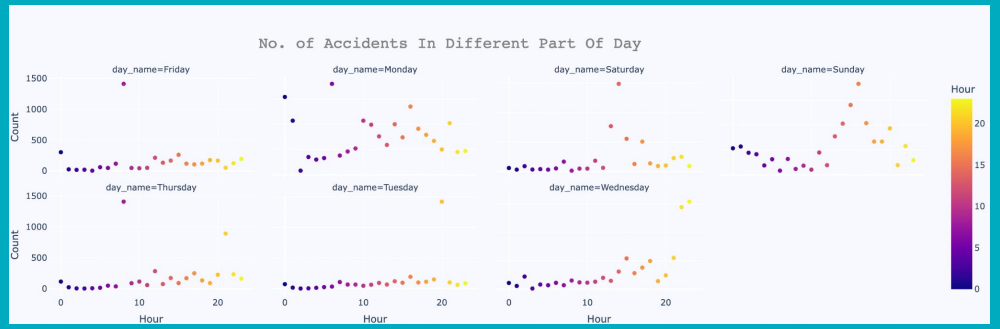

The charts indicate that most of the traffic accidents are at evening hours.

VEHICLE TYPE CODE Visualisation

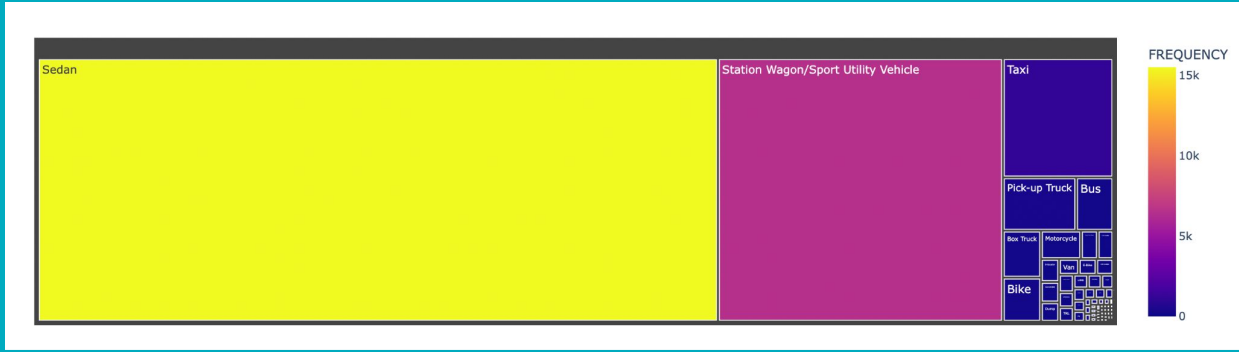

The treemap shows that Sedan is the most common type of vehicle in the accidents for VEHICLE TYPE CODE 1

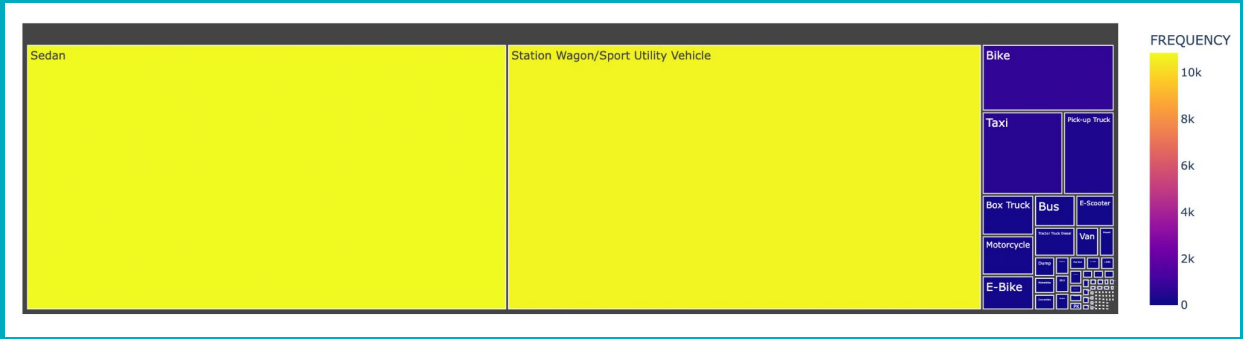

The treemap shows that Sedan is also the most common type of vehicle in the accidents for VEHICLE TYPE CODE 2

DATA MODELLING & MACHINE LEARNING

ONE-HOT ENCODING

PD.GET_DUMMIES

Water	Temperature
A	Hot
B	Cold
C	Warm
D	Cold

Dummy Variables

Water	Temperature	var_hot	var_warm	var_cold
A	Hot	1	0	0
B	Cold	0	0	1
C	Warm	0	1	0
D	Cold	1	0	0

Encoded DataFrame

	CRASH DATE	CRASH TIME	LATITUDE	LONGITUDE	NUMBER OF PERSONS INJURED	Hour	CONTRIBUTING FACTOR VEHICLE 1_Accelerator Defective	CONTRIBUTING FACTOR VEHICLE 1_Aggressive Driving/Road Rage	CONTRIBUTING FACTOR VEHICLE 1_Alcohol Involvement	CONTRIBUTING FACTOR VEHICLE 1_Animals Action	...	day_name_Sunday
0	2020-08-29	2022-11-09 14:00:00	40.704422	-73.792854	0	14	0	0	0	0	...	0
1	2020-08-29	2022-11-09 12:29:00	40.861862	-73.912820	2	12	0	0	0	0	...	0
3	2020-08-29	2022-11-09 19:00:00	40.839680	-73.929276	1	19	0	0	0	0	...	0
4	2020-08-29	2022-11-09 05:40:00	40.858190	-73.884350	0	5	0	0	0	0	...	0
9	2020-08-29	2022-11-09 15:00:00	40.669518	-73.911934	0	15	0	0	0	0	...	0
...	...	...	...	...	...	...	...	...	...	...	...	...
49528	2020-02-13	2022-11-09 08:25:00	40.665230	-73.931465	8	8	0	0	0	0	...	0
49529	2020-02-14	2022-11-09 08:40:00	40.854744	-73.923510	8	8	0	0	0	0	...	0
49530	2020-02-14	2022-11-09 08:40:00	40.854744	-73.923510	8	8	0	0	0	0	...	0
49533	2020-02-13	2022-11-09 08:25:00	40.665230	-73.931465	8	8	0	0	0	0	...	0
49534	2020-02-14	2022-11-09 08:40:00	40.854744	-73.923510	8	8	0	0	0	0	...	0

24536 rows × 246 columns

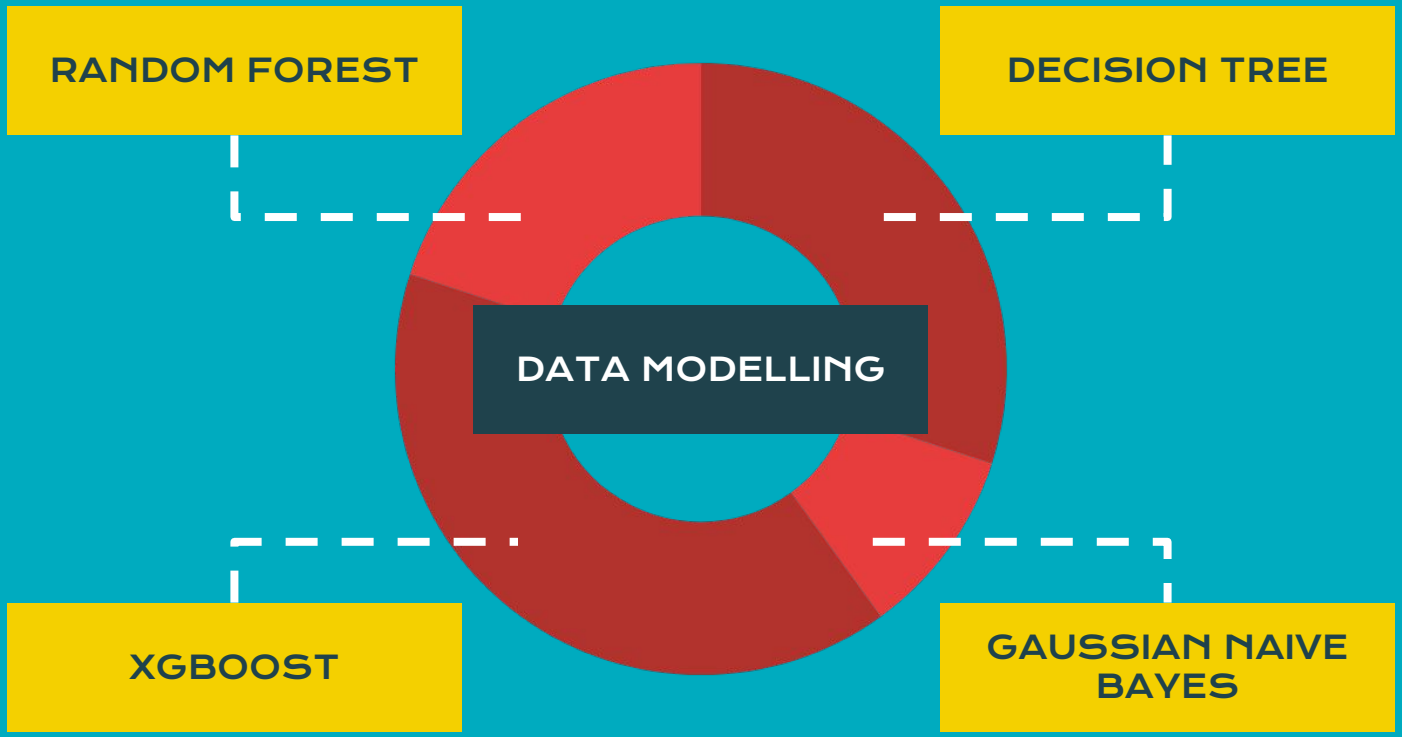
RANDOM FOREST
DECISION TREE
DATA MODELLING
XGBOOST
GAUSSIAN NAIVE BAYES

RANDOM FOREST CLASSIFICATION

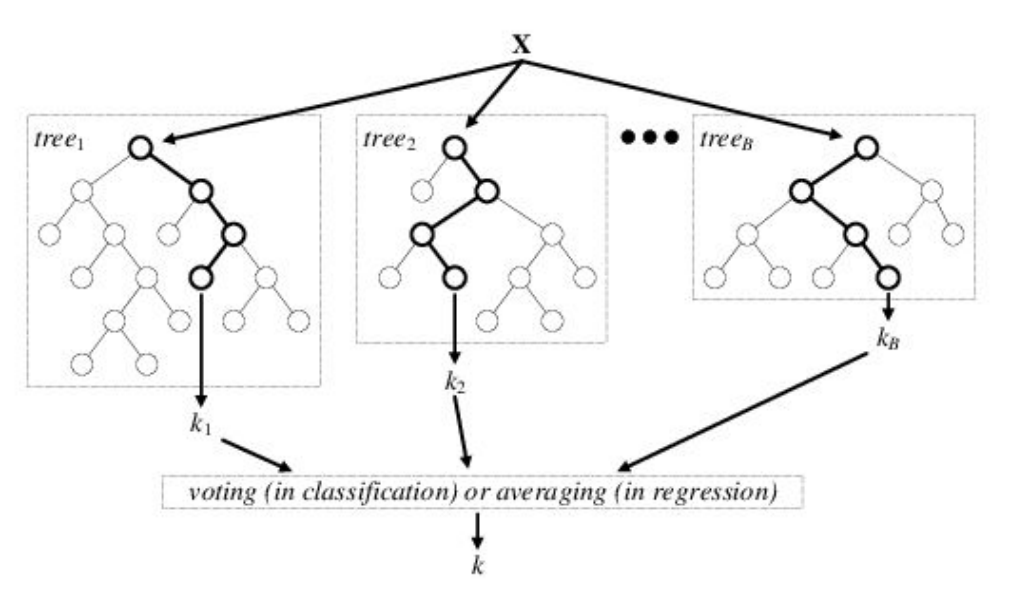

BOOTSTRAP AGGREGATION

Bagging is an ensemble algorithm that fits multiple models on different subsets of a training dataset, then combines the predictions from all models.

N_ESTIMATORS	100

We use 100 different decision trees for optimality.

RANDOM FOREST CLASSIFICATION

FEATURE IMPORTANCES

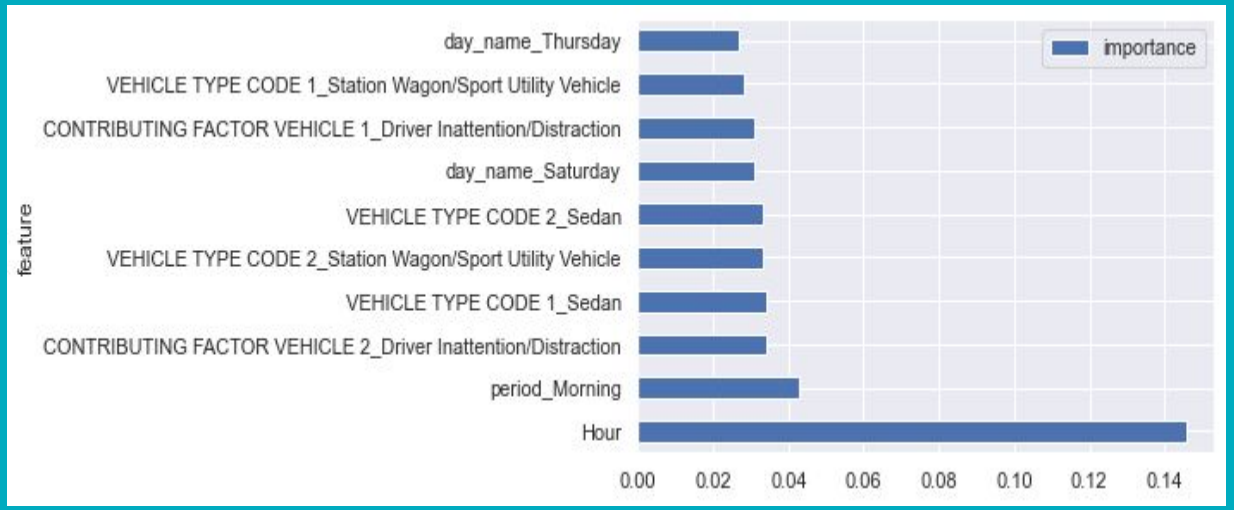

CONFUSION MATRIX

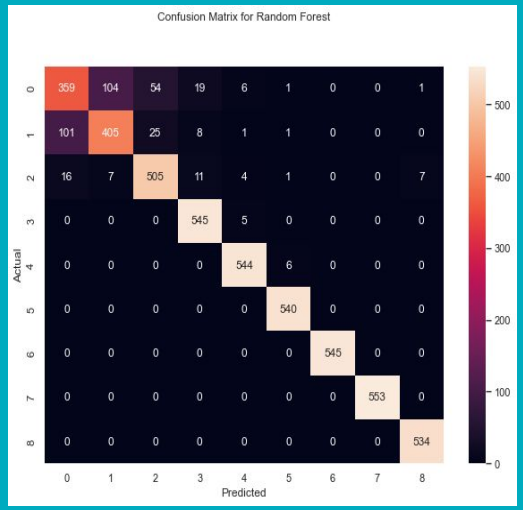

Accuracy	:	*0.9229828850855746*
Balanced Accuracy Score	:	*0.9227839528287662*
Precision	:	*0.9229828850855746*
Recall	:	*0.9229828850855746*

DECISION TREE CLASSIFICATION

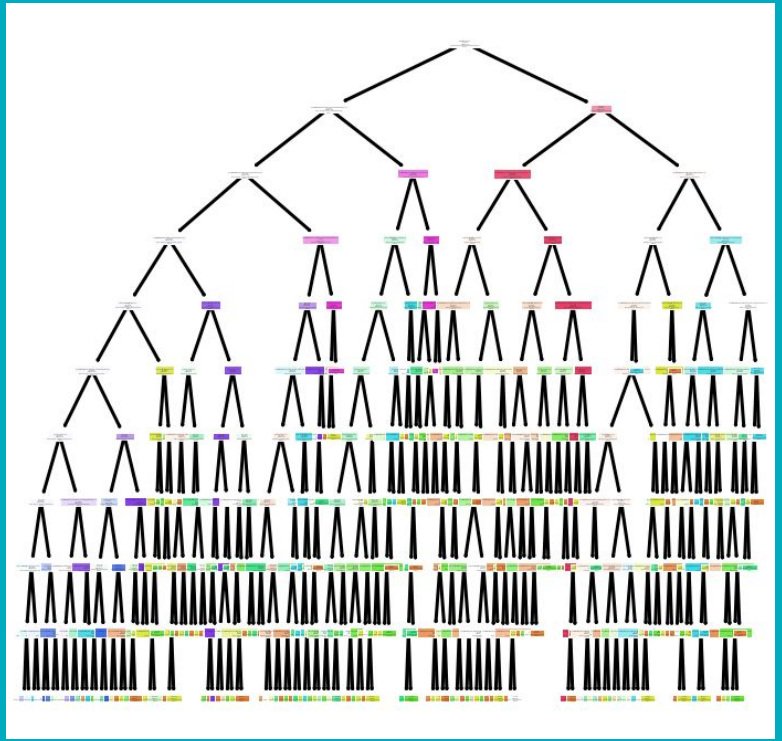

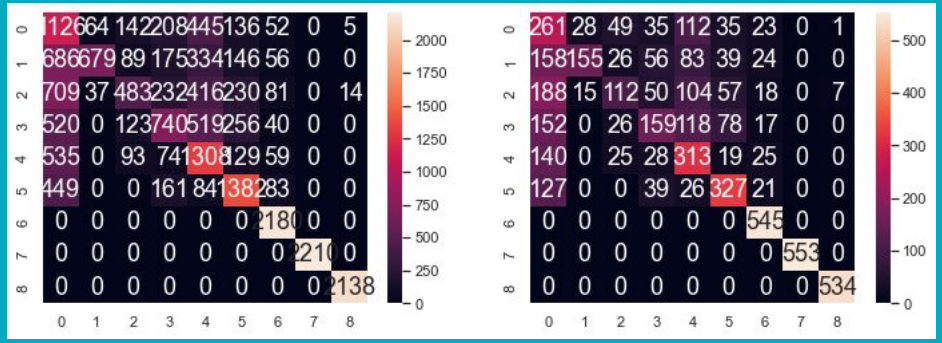

ACCURACY FOR TRAIN	0.6239046260
ACCURACY FOR TEST	0.6028932355

We need to optimise/tune the parameters!

Finding the Optimal Max_Depth using 5-Fold Cross Validation

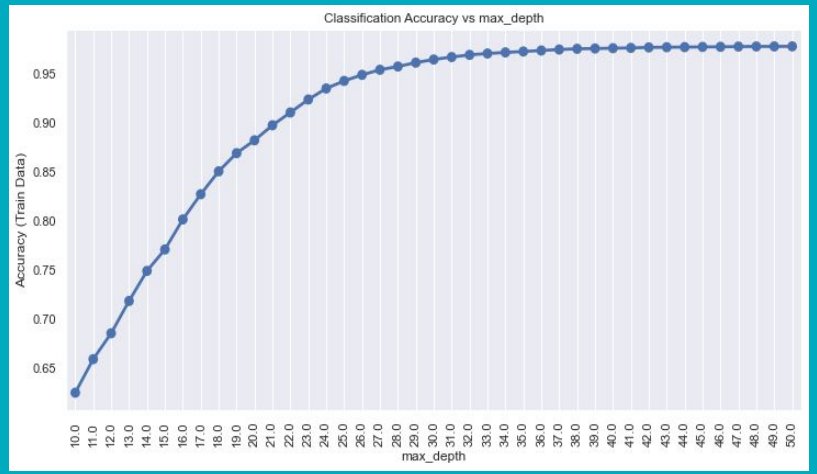

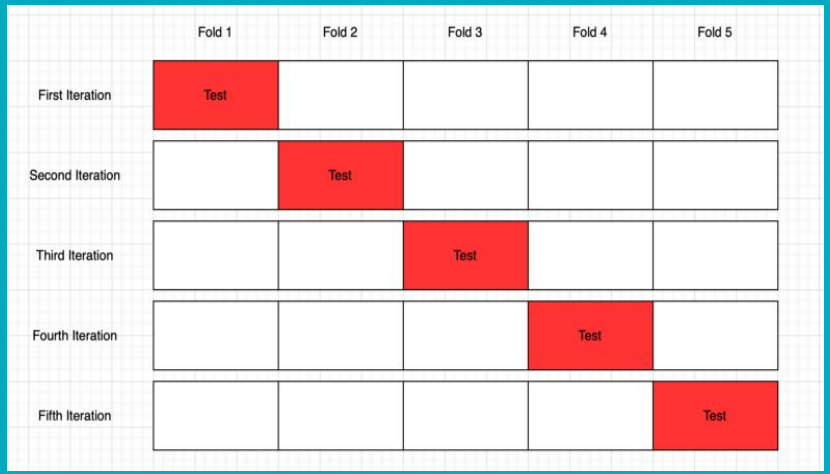

GridSearchCV

Grid search is a process that searches exhaustively through a manually specified subset of the hyperparameter space of the targeted algorithm.

StratifiedKFold

5

KFold divides all the samples in groups of samples. The fold left out is used for test. We use 5-fold cross validation.

DECISION TREE CLASSIFICATION

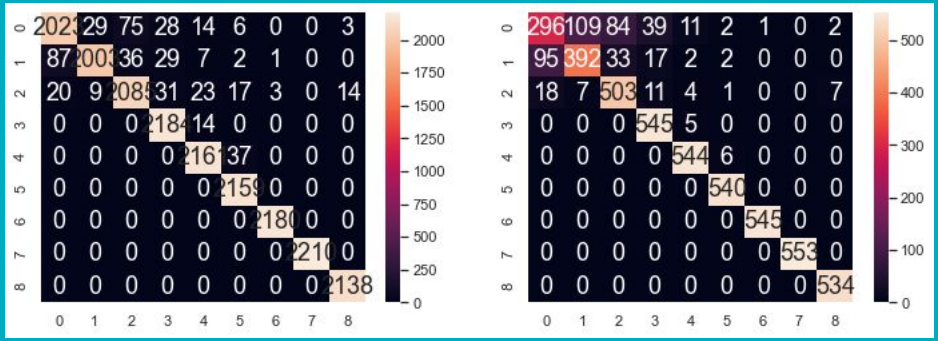

ACCURACY FOR TRAIN	0.97529040147
ACCURACY FOR TEST	0.907090645

The most optimal max_depth used is 39.

XGBOOST CLASSIFICATION

XGBoost, which stands for Extreme Gradient Boosting, is a scalable, distributed gradient-boosted decision tree (GBDT) machine learning library.

ONE TREE CREATED AT A TIME

ITERATION TO REDUCE ERRORS

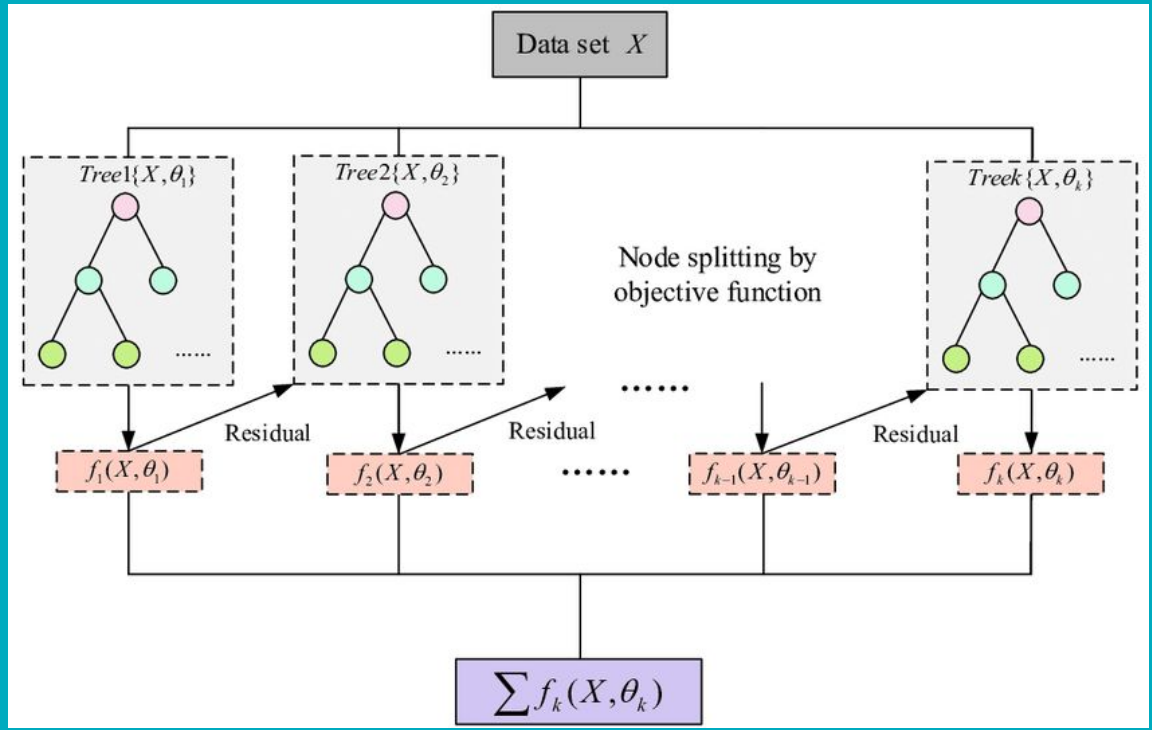

XGBOOST CLASSIFICATION

```
XGBClassifier(base_score=0.5, booster='gbtree', callbacks=None,
              colsample_bylevel=1, colsample_bynode=1, colsample_bytree=1,
              early_stopping_rounds=None, enable_categorical=False,
              eval_metric=None, feature_types=None, gamma=0, gpu_id=-1,
              grow_policy='depthwise', importance_type=None,
              interaction_constraints='', learning_rate=0.300000012,
              max_bin=256, max_cat_threshold=64, max_cat_to_onehot=4,
              max_delta_step=0, max_depth=6, max_leaves=0, min_child_weight=1,
              missing=nan, monotone_constraints='()', n_estimators=100,
              n_jobs=0, num_parallel_tree=1, objective='multi:softprob',
              predictor='auto', ...)
```

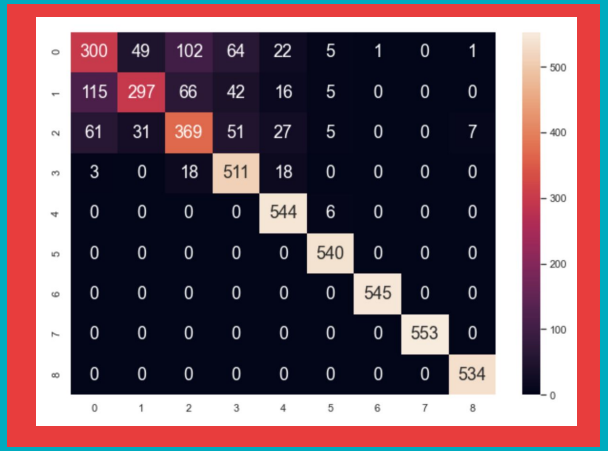

A **N_ESTIMATORS**

Based on the classification, the model requires to run **100** times to learn the data.

B **MAX_DEPTH**

Maxed depth we chose is as default, **6**, to prevent overfit of data

C **TEST ACCURACY**

Accuracy of model turned out positive with approximately **85.4%** success rate

GAUSSIAN NAIVE BAYES

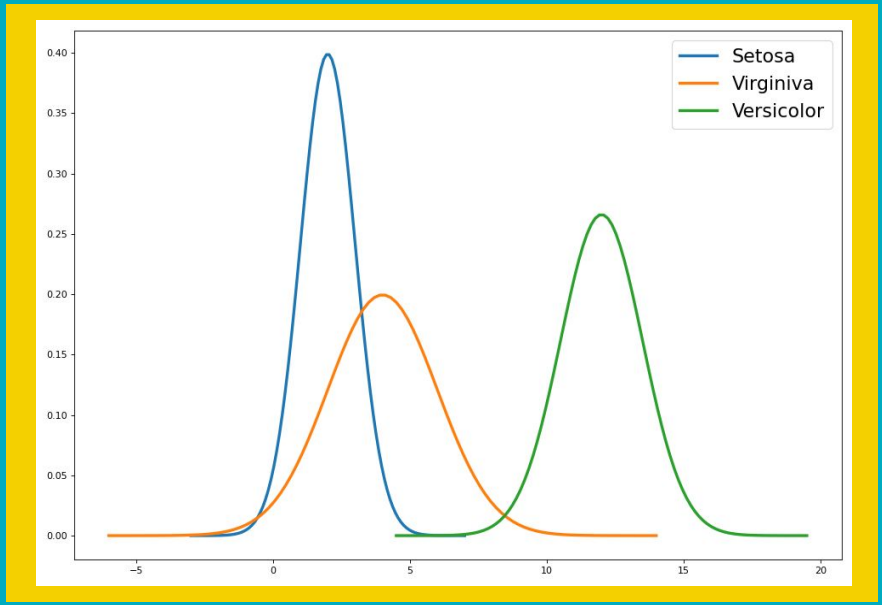

Naive Bayes is a classification algorithm for binary (two-class) and multi-class classification problems. The technique is easiest to understand when described using binary or categorical input values.

Assumes each class follows a Gaussian Distribution

Assumes that the features are independent

GAUSSIAN NAIVE BAYES

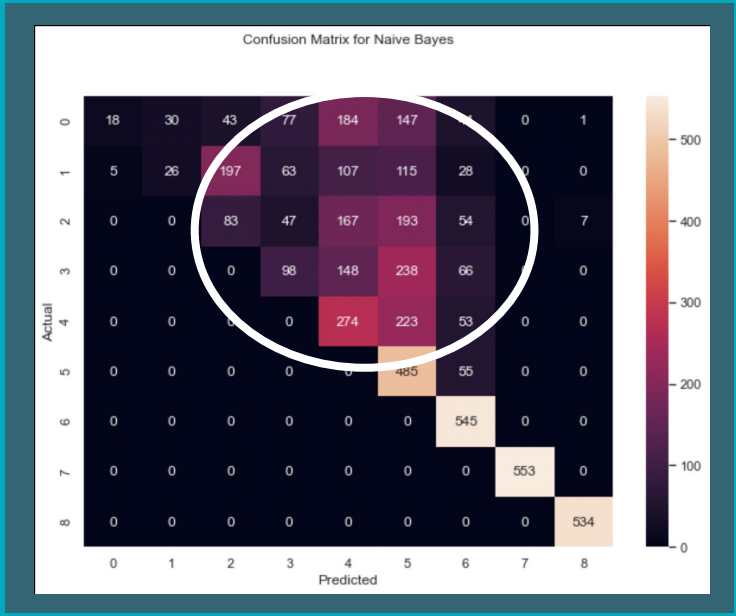

BALANCED ACCURACY **0.5340327087**

WHY IS IT SO LOW?

The Gaussian Naive Bayes assumes the dataset to be distributed normally.

In which our case, it is not.

CLOSING

CONCLUSION

Top 5 Most Important Features Predicting Number of Persons Injured

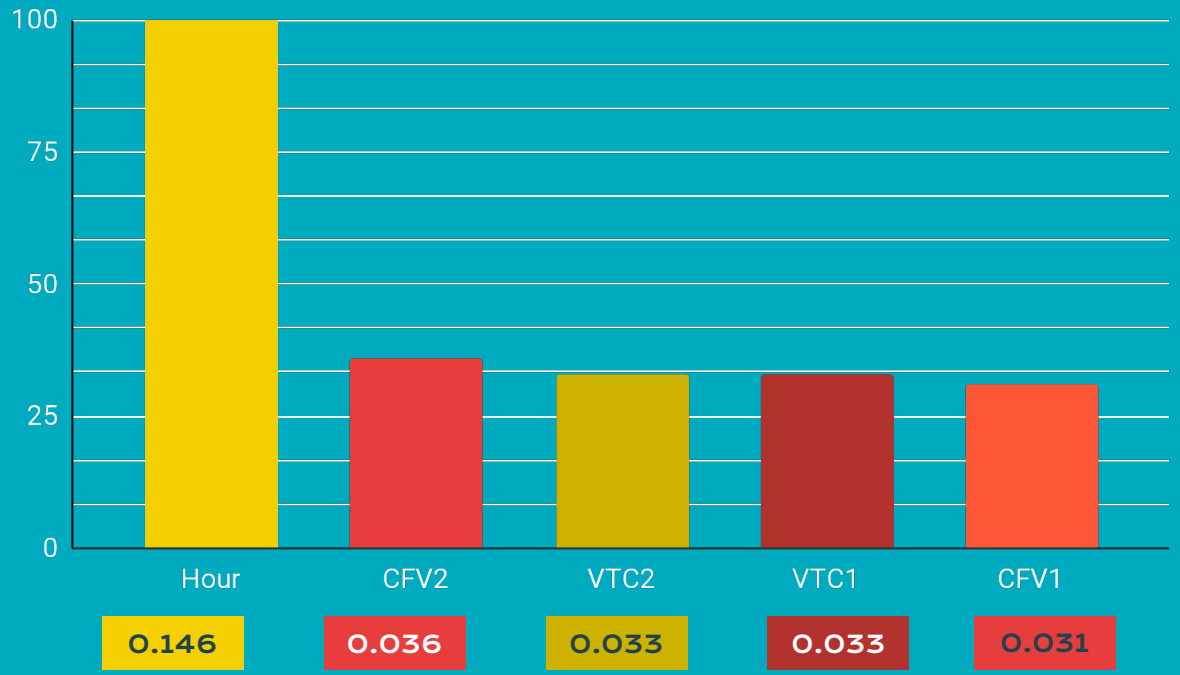

CONCLUSION

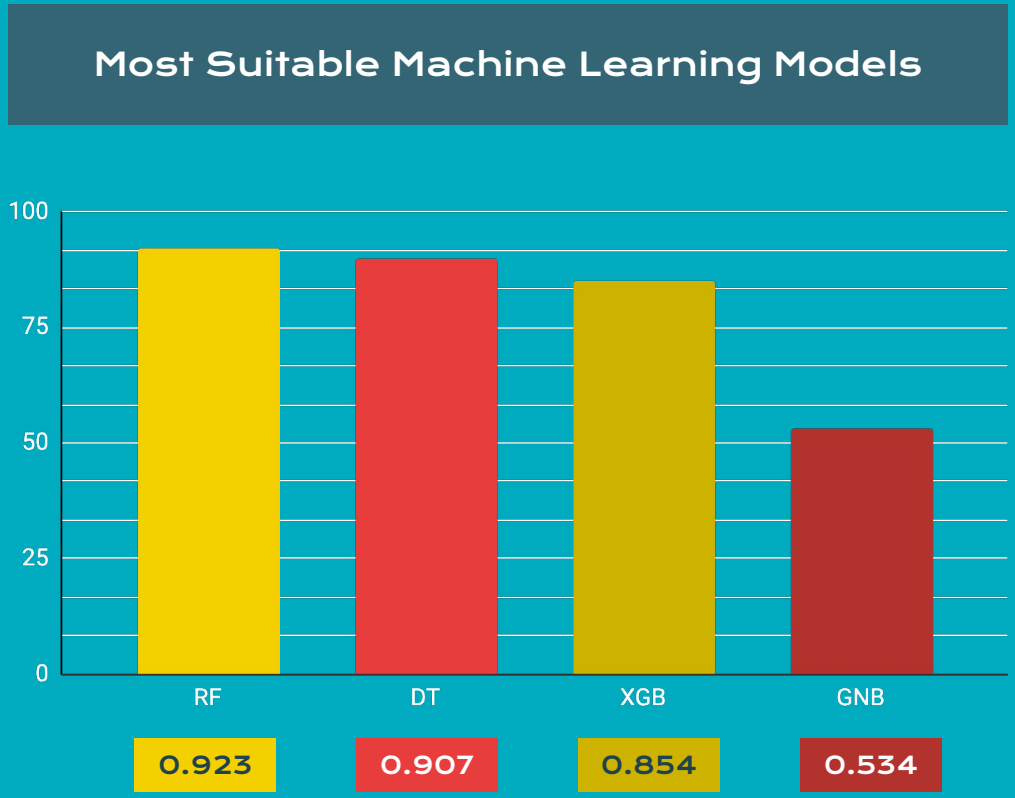
Most Suitable Machine Learning Models
100
75
50
25
0
RF
DT
XGB
GNB
0.923
0.907
0.854
0.534

Work Contribution

DATA PREPARATION

Francis & Matthew

DATA MODELLING & ML

Francis & Matthew

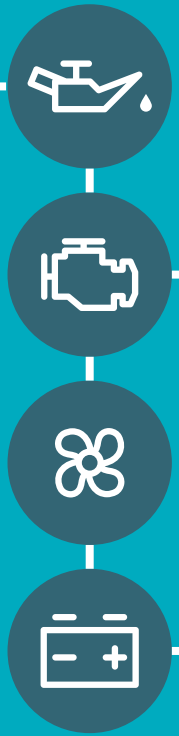

DATA VISUALISATION

Muh. Afiq & Nigel

CLOSING

Muh. Afiq & Nigel

THANK YOU

References

- https://www.sciencedirect.com/topics/mathematics/grid-search
- https://towardsdatascience.com/ensemble-methods-bagging-boosting-and-stacking-c9214a10a205
- https://towardsdatascience.com/gaussian-naive-bayes-4d2895d139a
- https://www.nvidia.com/en-us/glossary/data-science/xgboost/#:~:text=XGBoost%2C%20which%20stands%20for%20Extreme,%2C%20classification%2C%20and%20ranking%20problems.
- https://xgboost.readthedocs.io/en/stable/parameter.html